Ship.

Primary CI clears zero on the benefit side. Guardrails clean. Late-window stable.

TL;DR

SHIP. **Roll the change to 100% of users.**

Lift of +1.70pp on conversion_rate cleared the 2.0% MDE bar and all guardrails came back clean. Late-window effect held steady.

LIFT ON CONVERSION_RATE

+1.70pp

treatment vs control · +17.2% relative

95% CI ON LIFT
[+0.49, +2.91]pp

P-VALUE
0.0061

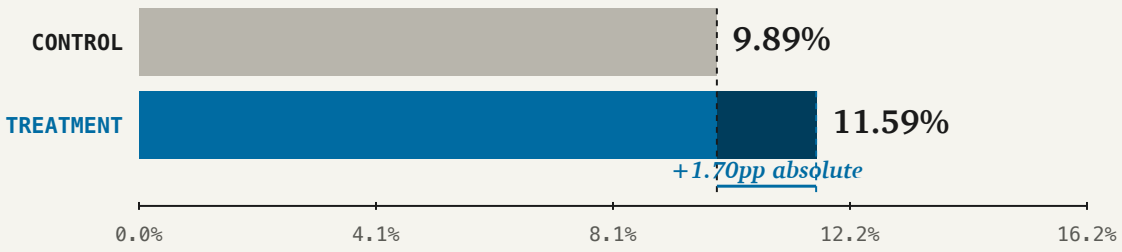
CONFIDENCE
highly likely positive

DECISION STEP
step 7

PRIMARY METRIC · CONVERSION_RATE

Treatment 11.59% vs Control 9.89%

Treatment and control rates on the same axis. The dark blue segment on the treatment bar is the absolute lift; the dotted line shows the lift relative to the control end.

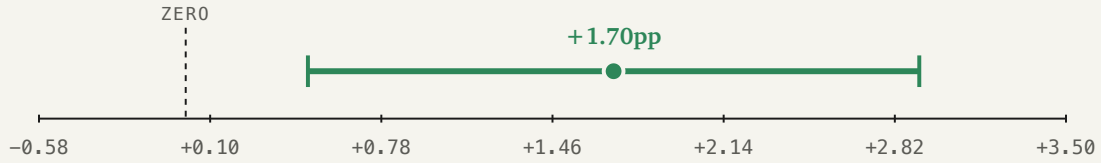


Per-arm conversion. The lift is +1.70pp absolute (+17.2% relative).

CONFIDENCE INTERVAL ON LIFT

95% CI on lift: (+0.49, +2.91)pp — clears zero on the benefit side.

The whole confidence interval lies above zero. We can rule out the null at $\alpha = 0.05$; the magnitude is what the verdict tree then weighs against the MDE.



The CI bar shows the range of plausible true lifts. The crossbars are the 95% bounds; the dot is the point estimate.

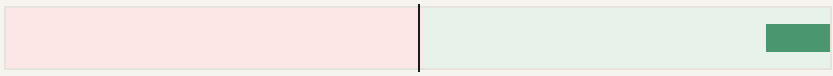
GUARDRAILS

Guardrail clear.

The pre-registered guardrails are revenue and (where wired) latency. None crossed their harm thresholds.

revenue_per_user

`threshold ≥ −5.0% · higher_is_better`



`90% CI [+12.66, +22.66]% · realized +17.66% ·` **CLEAR**

DECISION TREE · HOW THE VERDICT WAS REACHED

Walked 8 steps; step 7 fired the verdict.

Each row is a checkpoint in the verdict tree. Walking stops at the row marked FIRED; that step's logic is what produced the verdict above.

0	Required inputs present (UNVERIFIABLE guard)	`all non-null`
1	SRM passes (or override resolved)	`srm_pass = True`
2	Guardrails clear (no 90% CI on harm side)	`guard_lift = +17.66%`
3	Powered (n_observed ≥ n_required)	`n = 10,000 / 8,000`
4	Primary CI excludes 0	`CI [+0.49, +2.91] pp`
5	Primary CI not on harm side	`direction = higher_is_better`
6	Lift magnitude ≥ MDE/2	`|lift| = 1.699pp vs MDE/2 = 0.100pp`
7	Novelty stable (late_ratio ≥ 0.7) · **FIRED**	`late_ratio = —`

PROVENANCE · AUDITABLE CHAIN

Reproducible from sealed brief + warehouse fixture.

Every quantitative claim above resolves to a row in the warehouse, a function in `agentxp.stats.*`, or a step in `agentxp.interpret.tree`. The chain hash below covers the brief, every stat result, and the verdict tree's terminal step.

EXPERIMENT_ID	exp_2026q2_checkout_above_fold
CHAIN_HASH	b563b8074d71d49329d749325b4f50c984fe4383f4232842bc4375ac79f24ab4
RENDER_STATUS	VERIFIED
AGENTXP_VERSION	0.3.0-v3
GENERATED_AT	2026-06-05T15:41:53Z
PRIMARY_METRIC	conversion_rate · proportion_test
SRM_TEST	srm_chi_squared · threshold 0.0005
TREE_MODULE	agentxp.interpret.tree.walk_tree
REPRODUCE	python scratch/walk_experiment.py exp_2026q2_checkout_above_fold

---