EXP_2026Q2_RECS_V2 · VERDICT REPORT · ENGINEER

RECOMMENDATIONS · PERSONALIZATION · 2026-06-05 · 30,000 UNITS · 14 DAYS

Ship with caveat.

Real benefit-side lift, but smaller than half the pre-registered MDE.

TL;DR

SHIP WITH CAVEAT. **Ship — but the effect is smaller than the pre-registered MDE/2. Document and revisit.**

Lift of +0.73pp is real (CI excludes zero) but below half the pre-registered MDE. The effect is real but the business case may not clear the per-user cost.

LIFT ON ADD_TO_CART_RATE

+0.73pp

treatment vs control · +7.4% relative

95% CI ON LIFT

[+0.05, +1.41]pp

P-VALUE

0.0368

CONFIDENCE

highly likely positive

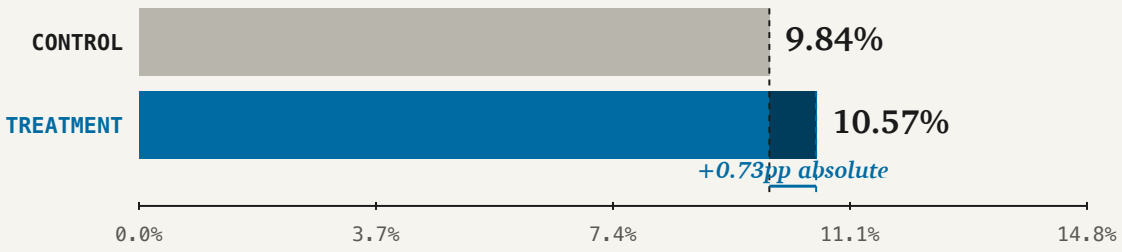
DECISION STEP

step 6

PRIMARY METRIC · ADD_TO_CART_RATE

Treatment 10.57% vs Control 9.84%

Treatment and control rates on the same axis. The dark blue segment on the treatment bar is the absolute lift; the dotted line shows the lift relative to the control end.

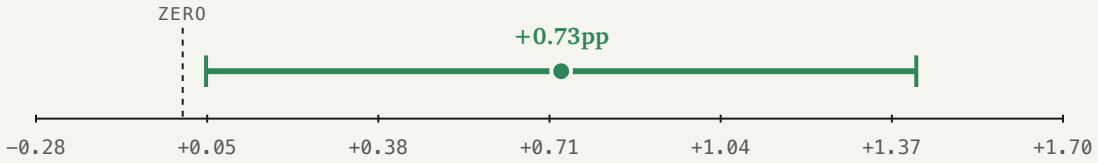


Per-arm conversion. The lift is +0.73pp absolute (+7.4% relative).

CONFIDENCE INTERVAL ON LIFT

95% CI on lift: (+0.05, +1.41)pp — clears zero on the benefit side.

The whole confidence interval lies above zero. We can rule out the null at α = 0.05; the magnitude is what the verdict tree then weighs against the MDE.



The CI bar shows the range of plausible true lifts. The crossbars are the 95% bounds; the dot is the point estimate.

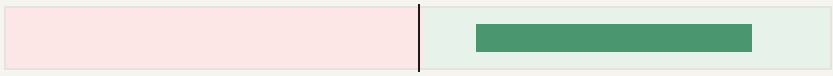
GUARDRAILS

Guardrail clear.

The pre-registered guardrails are revenue and (where wired) latency. None crossed their harm thresholds.

revenue_per_user

`threshold ≥ −5.0% · higher_is_better`



`90% CI [+2.13, +12.13]% · realized +7.13% ·` **CLEAR**

DECISION TREE · HOW THE VERDICT WAS REACHED

Walked 8 steps; step 6 fired the verdict.

Each row is a checkpoint in the verdict tree. Walking stops at the row marked FIRED; that step's logic is what produced the verdict above.

0	Required inputs present (UNVERIFIABLE guard)	`all non-null`
1	SRM passes (or override resolved)	`srm_pass = True`
2	Guardrails clear (no 90% CI on harm side)	`guard_lift = +7.13%`
3	Powered (n_observed ≥ n_required)	`n = 30,000 / 24,000`
4	Primary CI excludes 0	`CI [+0.05, +1.41] pp`
5	Primary CI not on harm side	`direction = higher_is_better`
6	Lift magnitude ≥ MDE/2 · **FIRED**	`|lift| = 0.730pp vs MDE/2 = 1.000pp`
7	Novelty stable (late_ratio ≥ 0.7)	`late_ratio = —`

PROVENANCE · AUDITABLE CHAIN

Reproducible from sealed brief + warehouse fixture.

Every quantitative claim above resolves to a row in the warehouse, a function in `agentxp.stats.*`, or a step in `agentxp.interpret.tree`. The chain hash below covers the brief, every stat result, and the verdict tree's terminal step.

EXPERIMENT_ID	exp_2026q2_recs_v2
CHAIN_HASH	31c7896bb5954e42a36a3d4dcff51468e0106812c89b561e84cdadd9abcb607e
RENDER_STATUS	VERIFIED
AGENTXP_VERSION	0.3.0-v3
GENERATED_AT	2026-06-05T15:42:27Z
PRIMARY_METRIC	add_to_cart_rate · proportion_test
SRM_TEST	srm_chi_squared · threshold 0.0005
TREE_MODULE	agentxp.interpret.tree.walk_tree
REPRODUCE	python scratch/walk_experiment.py exp_2026q2_recs_v2
